

Alejandro Sánchez Yalí

AI Full Stack Engineer



✉ asanchezyali@gmail.com | [in asanchezyali](#) | [asanchezyali](#) | [asanchezyali.com](#)

SUMMARY

AI Full Stack Engineer with 8+ years designing and shipping web applications end to end and 3+ years building LLM-powered products. Architect and lead engineer of three deployed systems – an AI WhatsApp support platform, an AI-assisted recruitment platform and a surgical supply chain system – spanning Python/FastAPI backends, Next.js/TypeScript frontends, PostgreSQL and cloud deployment. M.Sc. in Mathematics and 11 years teaching mathematics and physics, which shows up in design docs and code review.

TECHNICAL SKILLS

- **Languages:** Python (8+ years), TypeScript (6+ years), JavaScript (6+ years), SQL, Rust (personal projects)
- **Frontend:** React.js (7+ years), Next.js (6+ years), Tailwind CSS, Shadcn UI, Redux, Zustand, Zod, Effect, Three.js
- **Backend:** FastAPI, Django, Django REST Framework, Celery, Flask, NestJS, Node.js, Express, RESTful APIs, GraphQL, microservices
- **AI & LLMs:** LangChain, LangGraph, LiteLLM, RAG, MCP, OpenAI API, Whisper, ElevenLabs, PyTorch, TensorFlow
- **Databases & Infra:** PostgreSQL, Redis, MongoDB, MySQL, Firebase, Docker, AWS, GCP, GitHub Actions, CI/CD
- **Engineering:** Clean architecture, DDD, SOLID, TDD, Vitest, Jest, React Testing Library, code review

EXPERIENCE

Independent AI Consultant Oct 2026 – Present
Freelance Remote

- Lead engineer of Aluna, an AI recruitment platform for staffing agencies (see below)

AI Specialist Aug 2025 – Sep 2026
Lapzo Remote, Mexico

- Designed and built a Digital Professor that holds real-time conversations with learners and explains course content through interactive slides
- Owned the architecture connecting LLM services, ElevenLabs voice synthesis and the real-time layer built with Next.js, NestJS and Firebase Realtime Database
- Closed 57 tasks and 6 epics in the first nine months of 2026 with a single production bug, per the team's board
- Built the AI content-generation pipelines (LangChain, LangGraph, n8n), with Celery workers running generation in the background

Independent Software & AI Consultant Aug 2024 – Jul 2025
Freelance Remote

- Lead engineer of CREARIA's multi-tenant WhatsApp and email AI agent platform, in production with several client businesses including a real-estate agency: led a team of three across 5 FastAPI services and 2 Next.js dashboards on AWS ECS, with RAG over pgvector, multi-provider routing through LiteLLM, AI lead scoring and human escalation with full context; also led the AI coaching agent CREARIA delivers to one of its clients
- Delivered end-to-end web applications for business clients: React.js/Next.js/TypeScript on the front, Python/FastAPI/Django and Node.js/NestJS behind it, deployed on AWS with CI/CD and Stripe payments
- Ran technical design sessions with clients and made the technology decisions for each project; code reviewer for a client's development team

Full Stack Engineer Apr 2021 – Aug 2024
Monadical Remote, Canada/US

- Full-stack engineer on ZoHuddle, a video conferencing app for scientific conferences (Neuromatch) that matches participants by mutual interests: React front end and Django back end, with Django Channels (WebSockets) and Twilio Video
- Full-stack engineer on OddSlingers, an in-browser poker platform for desktop and mobile: React front end and Django back end on an event-stream architecture

- Built VC Tools, an LLM tool that retrieves and processes company information for investors (OpenAI, LangChain)
- Did code reviews on a distributed team across Canada, the US and Latin America

Full Stack Engineer
BCFort

Aug 2018 – Oct 2020
Medellín, Colombia

- Built blockchain dApp frontends (React, Web3.js) and analytics platforms (Django, Node.js) on AWS, GCP and IBM Cloud

Professor of Mathematics & Physics (part-time)
Universidad de Antioquia

Jan 2010 – Oct 2021
Medellín, Colombia

- Taught mathematics and physics; in the last year and a half also machine learning and data science with Python, NumPy, Pandas, TensorFlow and PyTorch
- Developed the AI/ML elective curriculum for the Mathematics Institute and mentored student research projects

FEATURED PROJECTS

- [Plixiq – AI Customer Support on WhatsApp](#)  – *Python, FastAPI, LiteLLM, PostgreSQL, Redis, Next.js, WebSocket*
 - Architect and lead engineer of a multi-tenant SaaS where AI agents answer routine customer inquiries, book appointments and escalate the rest to human operators with full conversation context
 - Modular monolith (14 components plus a shared kernel) with boundaries enforced in CI; LiteLLM layer with a primary and fallback provider per agent, so an outage becomes a fallback
- [VitaStock – Surgical Supply Chain for Clinics](#)  – *Python, FastAPI, Next.js, TypeScript, PostgreSQL, Docker, Railway*
 - Architect and lead engineer of a platform commissioned by a private surgical clinic in Colombia – 12 modules and 5 user roles designed to replace spreadsheet-based inventory and procurement; delivered, with rollout at the clinic pending
 - Surgery scheduling with supply baskets and real-time consumption tracking, lot-level traceability with expiry alerts, and true per-patient cost through a cardex
- [Aluna – AI Recruitment Platform for Staffing Agencies](#)  – *Next.js 15, FastAPI, LiteLLM, PostgreSQL, Inngest, Claude, Gemini, WhatsApp Cloud API*
 - Lead engineer of an AI recruitment platform for staffing agencies – 336 source files across 4 packages, 26 tables and 11 external providers – split into a Next.js 15 application owning all domain state and a stateless FastAPI service owning model calls and the WhatsApp channel
 - Durable Inngest analysis pipeline (9 checkpointed steps, 3 retries) where deduplication runs before quota enforcement so a re-analysed CV reuses its cached score instead of paying for another model call; one screening engine serves WhatsApp, web chat and email

OPEN SOURCE

- [talking-avatar-with-ai](#)  – 464 stars, 110 forks. Open-source conversational AI avatar: OpenAI GPT, Whisper, ElevenLabs, Rhubarb lip-sync, Three.js
- [Monadical-SAS/Morpheus](#)  – 19 pull requests. Open-source AI image generation platform: FastAPI, React, Stable Diffusion

TALKS & WRITING

- [Scaling Learning Models: Parallelism Strategies in JAX with Flax](#)  – PyCon Colombia, 2024
- **Engineering case studies:** [Plixiq](#)  · [CREARIA Agent](#)  · [Aluna](#)  – architecture and trade-offs behind each system (2026)

EDUCATION

M.Sc. in Mathematics
Universidad de Antioquia

Aug 2010 - Sep 2013
Medellín, Colombia

B.Sc. in Teaching Mathematics and Physics
Universidad de Antioquia

Jan 2004 - Oct 2009
Medellín, Colombia